



BT4014

Final Presentation

Movie Recommendation System

Ng Hong Yao (A0235009M)

Toh Hui Shan Alicia (A0204411B)

Vinessa Christabella (A0240431X)



Table of contents

01

Introduction

02

Dataset

03

ϵ -Decay

04




Contextual
Bandit

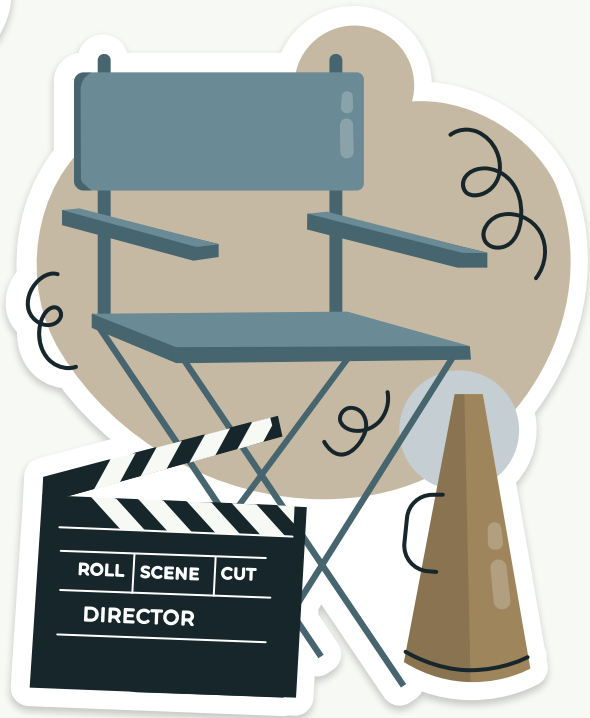
05

Limitations

06

Conclusion





01

Introduction



Background / Problem

- Online community where users can **rate** movies, **write** movie reviews and receive **movie recommendations**
- Problem of **under-contribution**
- **Inaccurate recommendation**, decrease in users' satisfaction & willingness to use the platform

movielens

Non-commercial, personalized movie recommendations.

[sign up now](#)

or [sign in](#)


recommendations

MovieLens helps you find movies you will like. Rate movies to build a custom taste profile, then MovieLens recommends other movies for you to watch.

The screenshot displays the MovieLens homepage. At the top, the 'movielens' logo is on an orange background. Below it, a navigation bar contains 'sign up now' and 'sign in' buttons. The main content area is white and features two primary sections: 'top picks' and 'recent releases'. The 'top picks' section is titled 'based on your ratings, MovieLens recommends these movies' and shows a row of movie cards including 'Band of Brothers', 'Casablanca', 'One Flew Over the Cuckoo's Nest', 'The Lives of Others', 'Sunset Boulevard', 'The Third Man', and 'Pulp Fiction'. Each card displays the movie title, year, duration, and a star rating. The 'recent releases' section is titled 'movies released in last 90 days that you haven't rated' and shows a row of movie cards including 'CartelFlix', 'Felixity', 'What If', 'Frank', 'Sin City: A Dame to Kill', and 'If I Stay'. Each card also displays the movie title, year, duration, and a star rating.

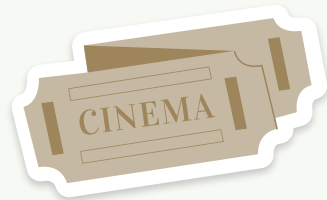
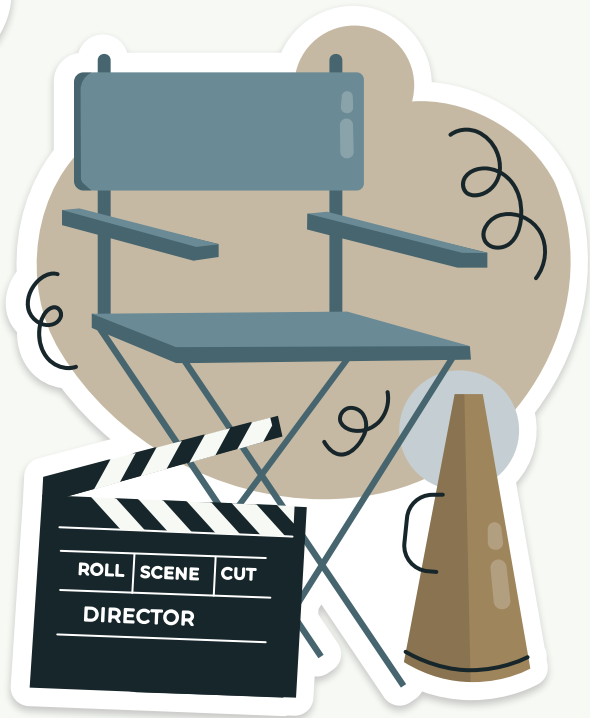
Proposed Solution

- Bandit Algorithms
 - **Epsilon decay** bandit algorithm combined with a content-based recommendation system
 - **Contextual bandit** algorithm which considers user features
- Goal: Increase user satisfaction by recommending high-quality movies



Content Based Recommendation System





02

Dataset



Dataset Description

- MovieLens dataset obtained from the grouplens website
- 100K variant of the data set, consisting of 100,000 ratings from 943 users on 1682 movies and each user has rated at least 20 movies
- 3 tables of interest:
 - **Users data:** Includes user ID with corresponding user demographics
 - **Movies data:** Includes information about movies based on their genres
 - **User Movie Ratings:** Contains users' ratings (1-5) for a movie



Dataset Preprocessing

Context Features


- Performing **data binning** on “age”
- 6 age group buckets: <20, 20-29, 30-39, 40-49, 51-60, >60
- Conduct **one-hot encoding** on categorical variables


Rewards

- Create **binary reward label**
- Ratings less than 4 are labeled as 0 and 1 otherwise

Filter

- Only include **30 movies** with most ratings on the website





03

Epsilon Decay





Epsilon Decay Modifications




Decay Function



Exploitation



Exploration



Decay Function

Original decay function

$$\epsilon = \frac{1}{\left(\frac{\text{number of play}}{\text{number of arms}} + 1 \right)}$$

Modified decay function

$$\epsilon = \frac{1}{\left(\frac{\text{number of movies watched by a user}}{\text{number of arms}} + 1 \right)}$$



Decay Function

$$\epsilon = \frac{1}{\left(\frac{\text{number of movies watched by a user}}{\text{number of arms}} + 1\right)}$$



User A
10 movies

Lower epsilon value
Exploit more



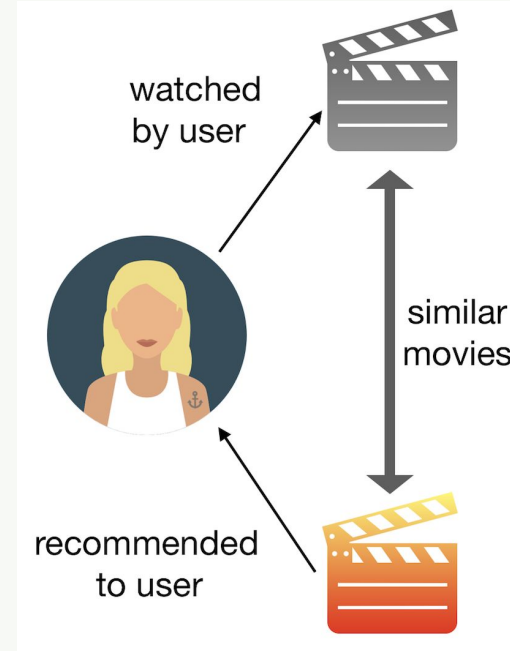
User B
2 movies

Higher epsilon value
Exploit less

Exploitation

Content Based Recommendation

suggests a movie which has similar characteristics to the movies a user has previously enjoyed



Exploitation

Content Based Recommendation

	movie_id	movie_title	release_date	genres
0	1	Toy Story (1995)	01-Jan-1995	Animation Children's Comedy
1	7	Twelve Monkeys (1995)	01-Jan-1995	Drama Sci-Fi
2	50	Star Wars (1977)	01-Jan-1977	Action Adventure Romance Sci-Fi War
3	56	Pulp Fiction (1994)	01-Jan-1994	Crime Drama
4	69	Forrest Gump (1994)	01-Jan-1994	Comedy Romance War
5	79	Fugitive, The (1993)	01-Jan-1993	Action Thriller
6	98	Silence of the Lambs, The (1991)	01-Jan-1991	Drama Thriller
7	100	Fargo (1996)	14-Feb-1997	Crime Drama Thriller

User A liked these movies (gave high ratings)

Recommend this to user A

Exploitation

Content Based Recommendation

	action	adventure	animation	children	comedy	crime	drama
0	0	0	1	1	1	0	0
1	0	0	0	0	0	0	1
2	1	1	0	0	0	0	0
3	0	0	0	0	0	1	1
4	0	0	0	0	1	0	0

How we achieve this:

- **Prepare data:** Matrix of token counts for movie genres
- **Algorithm:** Nearest neighbour algorithm with cosine distance metric
 - most similar to the history of movies that the user enjoyed before
- **Filter:** the user has not watched that movie before

*if no result, then recommend overall best movie so far

Exploration

Recommend recent movie

Method

Pick a random movie from 'k' most recent movies among the arms

Rationale

Recent movies are more likely to be appealing and relevant to viewers

Hyperparameter

$k \rightarrow$ controls the number of top recent movies selected for exploration

We explored 3 values of k (5, 15, 25)



Epsilon Decay Algorithm

Implementation:

- A class object for the content based recommendation system (called Recommender) that is used for exploitation part

```
class Recommender:
```

- For each k value, a function that contains epsilon decay algorithm



```
def epsilon_decay_k_5(user_history, history):
```

```
def epsilon_decay_k_15(user_history, history):
```

```
def epsilon_decay_k_25(user_history, history):
```

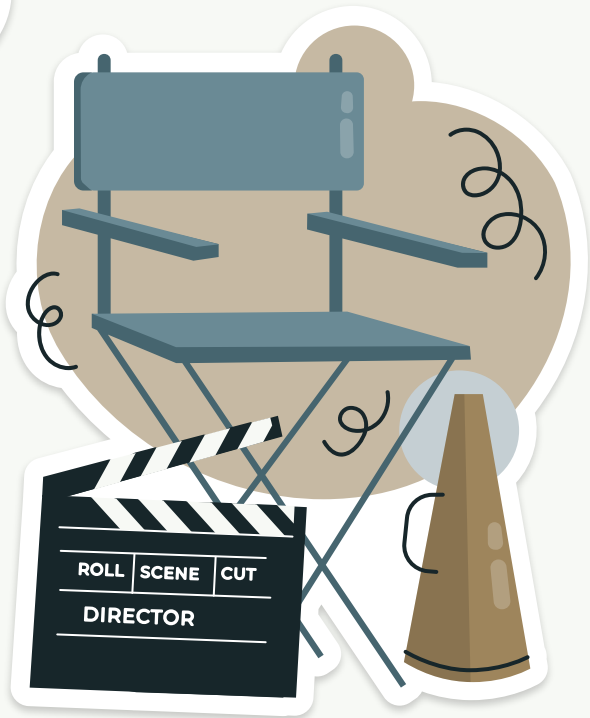

- Bootstrap resampling using 3 bootstrap samples

Epsilon Decay Algorithm



Cumulative reward over time:

- **k = 25** has **highest** cumulative rewards
- **k = 15** has the **lowest** cumulative rewards
- **k = 5** is higher than **k = 15**, lower k doesn't mean lower cumulative rewards



04

Contextual Bandit




Contextual Bandit

- Additional **context vector** is used during learning
- Information such as **user's demographics** (age, gender, occupation)
- Ability to make different recommendations to different groups of users with different tastes & preferences

Hyperparameter

- α , controls balance between exploration & exploitation
- **Higher** value of α , wider the confidence bound, **greater emphasis** on **exploration**



Contextual Bandit




Linear Upper
Confidence Bound

Disjoint




Linear Upper
Confidence Bound

Hybrid



LinUCB Disjoint

- Arms (movies) are **distinct** from each other
- Each time step:
 - Select arm with **greatest** UCB value
 - Observe reward for selected arm only
- Strength:
 - Ability to **explore** arms with potential but uncertain returns
 - While **exploiting** arms that are known to have high returns





LinUCB Disjoint

3 Objects For Implementation:



- A class object to represent a LinUCB disjoint arm
- A class object for the policy with the specified number of LinUCB disjoint arms
- A function that implements bootstrap replay using the LinUCB policy created above

```
class linucb_disjoint_arm():
```

```
class linucb_disjoint_policy():
```

```
def disjoint_bootstrap_replay(K_arms, d, alpha, top_movies_index, bootstrap_resample):
```

LinUCB Disjoint



Time Step < 5000:


- $\alpha = 1.0$ has **slowest** increase in cumulative rewards due to greater emphasis on **exploring arms with high uncertainty**

Time Step > 5000:

- Cumulative reward of $\alpha = 1.0$ surpass that of $\alpha = 0.25$
- $\alpha = 0.5$ has **greatest** amount of cumulative rewards

LinUCB Hybrid

- Arms (movies) are **not mutually exclusive** in properties
- Each time step:
 - Select arm with **greatest** UCB value
 - Observe reward for selected arm & **similar arms**
- Strength:
 - We account for **shared features** between arms
 - Reward payoff is a linear function of both non-shared and shared components





LinUCB Hybrid

3 Objects For Implementation:

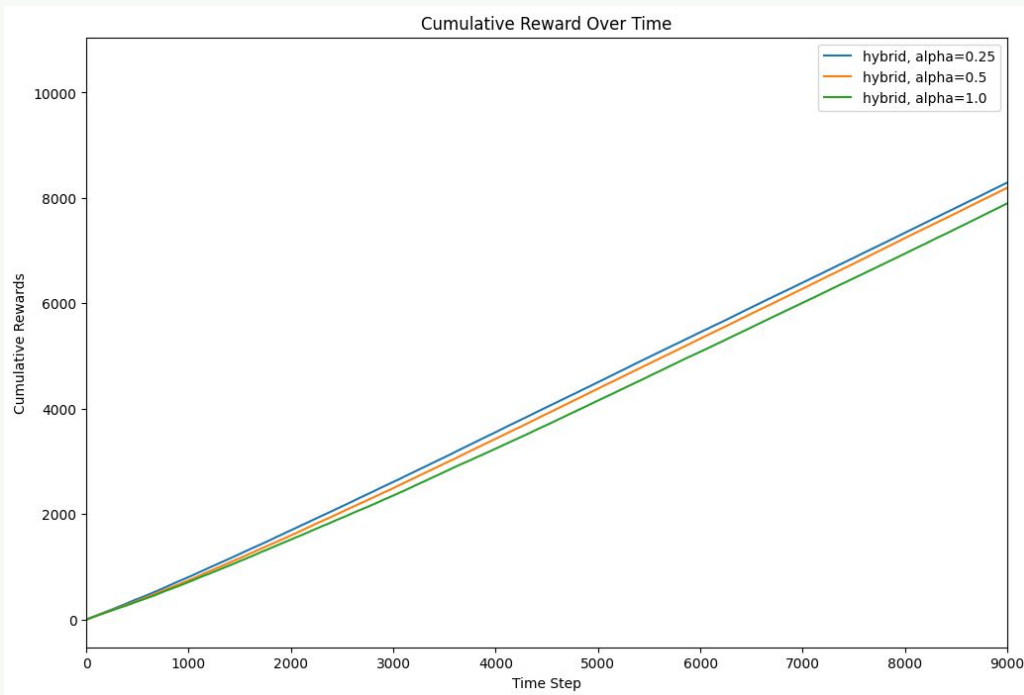

- A class object to represent a LinUCB hybrid arm
- A class object for the policy with the specified number of LinUCB hybrid arms
- A function that implements bootstrap replay using the LinUCB policy created above

```
class linucb_hybrid_arm():
```

```
class linucb_hybrid_policy():
```



```
def hybrid_bootstrap_replay(K_arms, d, k, alpha, top_movies_index, top_movies_features, bootstrap_resample):
```

LinUCB Hybrid




- $\alpha = 0.25$ has **greatest** amount of cumulative rewards at all time step
- Lower weightage of exploration
- Significantly **slower performance** for $\alpha = 1.0$

LinUCB Disjoint vs Hybrid



- **Hybrid** models had **greater** cumulative rewards than disjoint models
- Advantage: utilizes shared features of the arms to help model the reward function for arms with similar features




05

Limitations



Limitations



Overall project



- Due to resource constraints, we **filtered** the dataset to include only **a subset** of the available movies.

ϵ -Decay

- We only experimented with 3 **k values** (5, 15, 25) but other values could work better in the exploration phase
- We employed content-based recommender system as part of exploitation phase. There could be **other recommender system to be integrated.**

Contextual Bandit

- We only experimented with 3 **α values** (0.25, 0.5, 1.0) but other values could work better
- We could explore more **variations of contextual information** such as creating interaction terms from the existing covariates

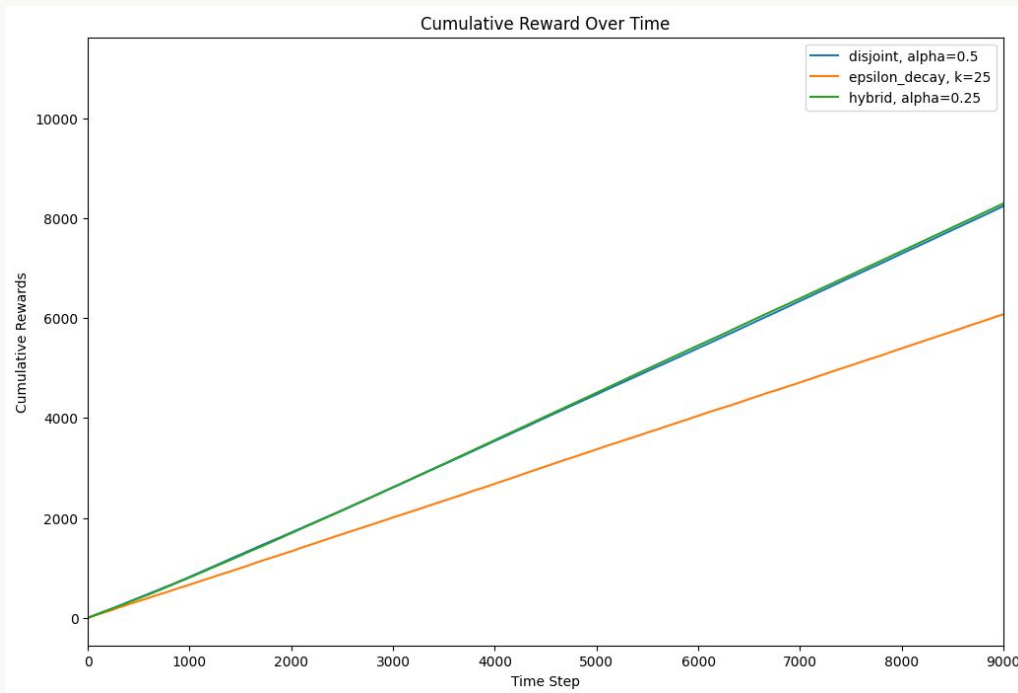



06


Conclusion



Conclusion



- **Best performing** model is LinUCB Hybrid, $\alpha = 0.25$ with greatest cumulative reward of 8293.33
- Epsilon decay, $k = 25$ performed **significantly worse** than LinUCB with cumulative reward of 6075.67



Thank You